

Introductory Concepts



Functionality comes with complexity?

- Every piece of computer hardware different
 - Different CPU - Pentium, PowerPC, ARM, MIPS
 - Different amounts of memory, disk, ...
 - Different types of devices
 - Mice, Keyboards, Sensors, Cameras, Fingerprint readers
 - Different networking environment
 - Cable, DSL, Wireless, Firewalls,...



Key Questions

- Questions:
 - Does the programmer need to write a single program that performs many independent activities?
 - Does every program have to be altered for every piece of hardware?
 - Does a faulty program crash everything?
 - Does every program have access to all hardware?



The OS paradigm

Two main functions:

Manage physical resources:

- It drives various devices
 - Eg: CPU, memory, disks, networks, displays, cameras, etc
- Efficiently, reliably, tolerating and masking failures, etc

Provide an execution environment to the applications running on the computer (programs like Word)

- Provide virtual resources and interfaces
 - Eg: files, directories, users, threads, processes, etc
- Simplify programming through high-level abstractions
- Provide users with a stable environment, mask failures



System design is Complex

- OS are a class of exceptionally complex systems
 - They are large, parallel, very expensive, not understood
 - OS implementation : years, large groups of people
 - Complex systems are the most interesting:
 - Internet, air traffic control, weather forecast, etc.
- How to deal with this complexity?
 - Abstractions and layering
 - Goal: systems trusted with sensitive data and critical roles



What are the issues in OS / System Design?

- Structure: how is an operating system organized ?
- Sharing: how are resources shared among users ?
- Naming: how are resources named by users or programs ?
- Protection: how is one user/program protected from another ?
- Security: how to authenticate, control access, secure privacy ?
- Performance: why is it so slow ?
- Reliability and fault tolerance: how do we deal with failures ?
- Extensibility: how do we add new features ?



What are the issues in OS / System Design? ...cont'd

- Communication: how can we exchange information ?
- Concurrency: how are parallel activities created and controlled ?
- Scale, growth: what happens as demands or resources increase ?
- Persistence: how can data outlast processes that created them
- Compatibility: can we ever do anything new ?
- Distribution: accessing the world of information
- Accounting: who pays bills, and how to control resource usage

Why study Operating Systems?

- Operating systems are a maturing field
 - Most people use a handful of mature OSes
 - Hard to get people to switch operating systems
 - Hard to have impact with a new OS
- High-performance servers are an OS issue
 - Face many of the same issues as OSes
- Resource consumption is an OS issue
 - Battery life, radio spectrum, etc.
- Security is an OS issue
 - Hard to achieve security without a solid foundation
- New “smart” devices need new OSes

Big Data, Networks, large-scale software systems, whatever --- Key to its success is Operating System Design Principles, its working and optimal use of its policies.

What is an OS?

- Tool to make programmer's job easy
- Resource allocator / manager
 - Must be fair; not partial to any process, specially for process in the same class
 - Must discriminate between different class of jobs with different service requirements
 - Do the above efficiently
 - Within the constraints of fairness and efficiency, an OS should attempt to maximize throughput, minimize response time, and accommodate as many users as possible
- Control program
- Tool to facilitate efficient operation of a computer system
- Virtual machine that is easier to understand and program
 - Encapsulates the complexities of hardware

OS as Resource Manager

- A computer is a set of resources for the movement, storage, and processing of data.
 - Data could be associated with programs, devices, and even the underlying hardware.
- The OS is responsible for managing these resources.

OS as Service Provider

- Program development
 - e.g., editors and debuggers
- Program execution
- Access I/O devices
- Controlled access to files
- System access for shared systems
- Error detection and response
 - e.g., memory error, device failure, division by zero
- Accounting for resources and performance monitoring


How about Operating System as a Service? - Think about it!
We are in the age of Cloud Computing – How OS changes?

Operating System

- A program that controls the execution of application programs
- An interface between applications and hardware
- Main objectives of an OS:
 - Convenience
 - Efficiency and Fairness
 - Ability to evolve


Uniprogramming

- Processor must wait for I/O instruction to complete before proceeding



Multiprogramming

- When one job needs to wait for I/O, the processor can switch to the other job



(c) Multiprogramming with three programs

Idea of degree of multiprogramming – solution to avoiding CPU being idle
But what is the effect on Response time?

Time Sharing Systems

- Using multiprogramming to handle multiple interactive jobs
- Processor's time is shared among multiple users

OS Concepts

- Kernel
 - Permanently resides in the main memory
 - Controls the execution of processes by allowing their creation, termination or suspension, and commn.
 - Schedules processes fairly for execution on the CPU
 - Allocates main memory for an executing process
 - File system maintenance
 - Allows processes controlled access to peripheral devices such as terminals, tape drives, disk drives, and network devices.
 - Design options – Monolithic kernels, Layered, Micro kernels.

- Kernel

- Permanently resident in the main memory
- Controls the execution of processes by allowing their creation, termination or suspension, and commn.
- Schedules processes fairly for execution on the CPU
- Allocates main memory for an executing process
- File system maintenance
- Allows processes controlled access to peripheral devices such as terminals, tape drives, disk drives, and network devices.
- Design options – Monolithic kernels, Layered, Micro kernels.

- Processes share the CPU in a time-shared manner
 - CPU executes a process
 - Kernel suspends it when its time quantum elapses
 - Kernel schedules another process to execute
 - Kernel later reschedules the suspended process

OS Concepts

- Kernel

Allows processes to share portions of their address space under certain conditions, but protects the private address space of a process from outside tampering

If system runs low on free memory, the kernel frees memory by writing a process temporarily to secondary memory, or swap device

If kernel writes entire processes to a swap device, then the implementation is called a swapping system; if it writes pages of memory to a swap device, it is called a paging system.

Coordinates with the machine hardware to set up a virtual to physical address that maps the compiler generated addresses to their physical addresses

- Allocates main memory for an executing process
- File system maintenance
- Allows processes controlled access to peripheral devices such as terminals, tape drives, disk drives, and network devices.
- Design options – Monolithic kernels, Layered, Micro kernels.

OS Concepts

- Kernel

- Permanently resides in the main memory
- Controls the execution of processes, process creation, termination or suspension
- Schedules processes fairly for execution
- Allocates main memory for processes
- File system maintenance
 - Allocates secondary memory for efficient storage and retrieval of user data
 - Allocates secondary storage for user files
 - Reclaims unused storage
 - Structures the file system in a well understood manner
 - Protects user files from illegal access
- Allows processes controlled access to peripheral devices such as terminals, tape drives, disk drives, and network devices.
- Design options – Monolithic kernels, Layered, Micro kernels.

OS Concepts - Program

- Program
 - Collection of source code instructions and any associated data kept in a disk file
 - Source program, or a human-readable text file.
 - Machine language translation of the source program, or object file
 - The file is marked as executable.
 - File contents are arranged according to rules established by the kernel
 - Executable program, complete code output by linker / loader, with input from libraries

Example of Program Execution




Figure 1.4 Example of Program Execution
(contents of memory and registers in hexadecimal)

Instruction Execution

- A program consists of a set of instructions stored in memory
- Two steps
 - Processor reads (fetches) instructions from memory
 - Processor executes each instruction




Figure 1.2 Basic Instruction Cycle

OS Concepts

- Memory (Volatility)
 - Memory hierarchy based on storage capacity, speed, and cost
 - Higher the storage capacity, lesser the speed, and lesser the cost
 - Different memory levels, in decreasing cost per byte of storage
 - Use hierarchical memory to transfer data from lower memory to higher memory to be executed
 - Locality of reference
 - Most of the references in the memory are clustered and move from one cluster to the next
 - Cache memory
 - Use of very fast memory (a few kilobytes) designated to contain data for fast access by the CPU
 - Virtual memory or extension of main memory
 - Disk cache
 - Data-intensive applications (generally rotational speed and seek time)

OS Concepts

- Memory (Volatility)

- Memory hierarchy
- Higher the storage capacity, slower the access time
- Different memory technologies
 - Registers
 - Cache memory
 - Main memory
 - Magnetic disk
 - Magnetic tape/Optical disk
- Use hierarchical memory to transfer data from lower memory to higher memory to be executed
- Locality of reference
 - Most of the references in the memory are clustered and move from one cluster to the next
- Cache memory
 - Use of very fast memory (a few kilobytes) designated to contain data for fast access by the CPU
- Virtual memory or extension of main memory
- Disk cache
 - Data-intensive applications (generally rotational speed and seek time)

Registers	Few bytes	Almost CPU speed
Cache memory	Few kilobytes	Nanoseconds
Main memory	Megabytes	Microseconds
Magnetic disk	Gigabytes	Milliseconds
Magnetic tape/Optical disk	No limit	Offline storage

OS Concepts

- Memory (Volatility)

- Memory

A disk cache is a mechanism for improving the time it takes to read from or write to a hard disk.

- High

Today, the disk cache is usually included as part of the hard disk.

- Different

(A disk cache is RAM built into your hard disk)

(Disk specifications : on-board cache)

- Use

OR

A disk cache can also be a specified portion of system RAM (disk buffer).

- Local

Disk cache holds data that has recently been read and, in some cases, adjacent data areas that are likely to be accessed next.

-

Write caching (accumulated) is also provided with some disk caches.

- Cache in

- Use of very small memory (a few kilobytes) designated to contain data for fast access to CPU

- Virtual memory or extension of main memory

- Disk cache

- Data-intensive applications (generally rotational speed and seek time)

OS Concepts – System Calls

Interface between user program and operating system - Set of extended instructions provided by the operating system

- Applied to various software objects like processes and files
- Invoked by user programs to communicate with the kernel and request services
- Access routines in the kernel that do the work
- Library procedure corresponding to each system call

Why modes?
Privileges?
- indicates power for
System control

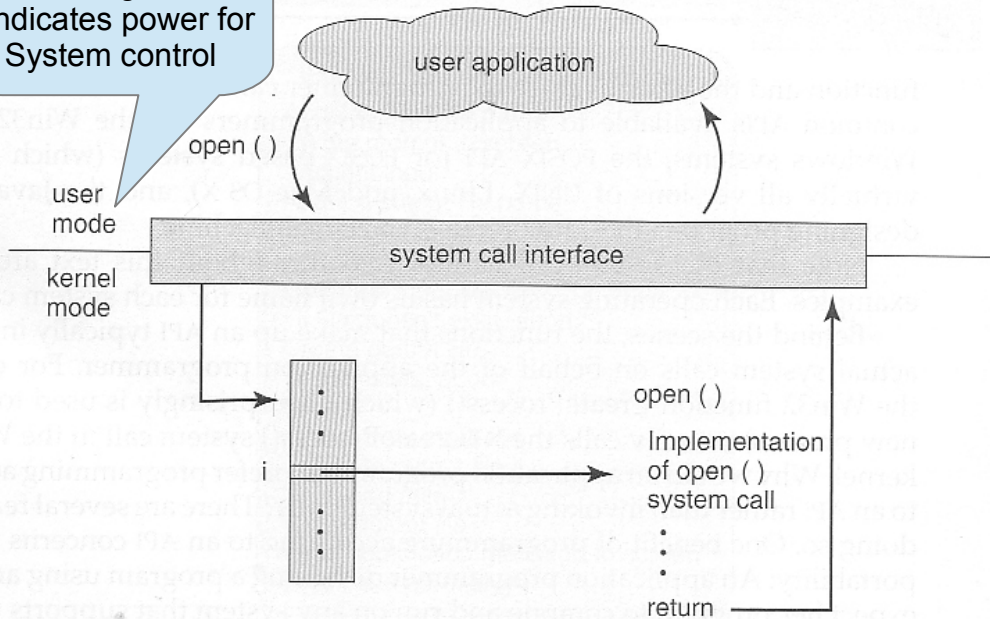


Figure 2.3 The handling of a user application invoking the `open()` system call.

EXAMPLE OF STANDARD C LIBRARY

The standard C library provides a portion of the system-call interface for many versions of UNIX and Linux. As an example, let's assume a C program invokes the `printf()` statement. The C library intercepts this call and invokes the necessary system call(s) in the operating system—in this instance, the `write()` system call. The C library takes the value returned by `write()` and passes it back to the user program. This is shown in Figure 2.6.




Figure 2.6 C library handling of `write()`.

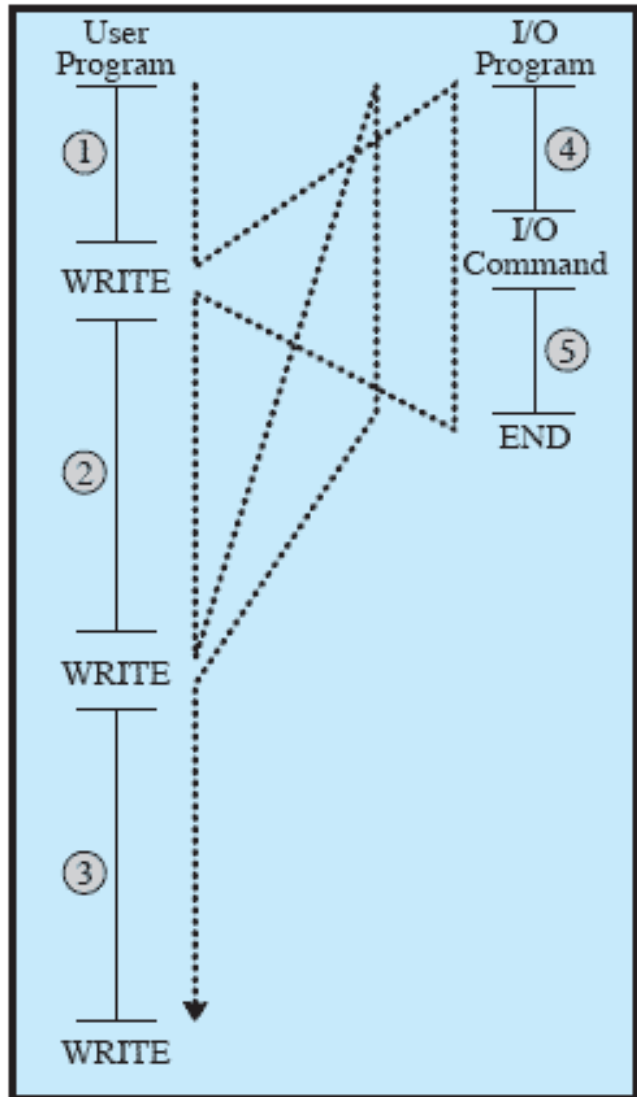
Working of a System Call – How?

Interrupts

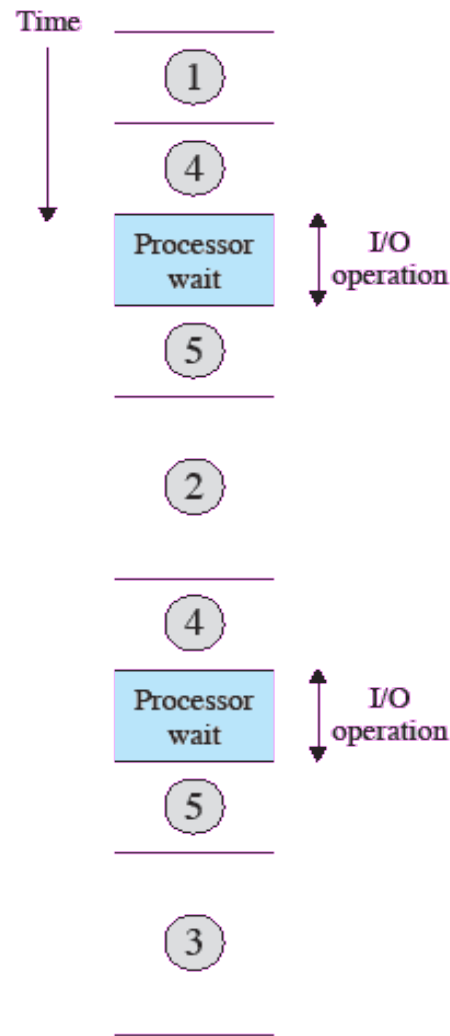
- Interrupt the normal sequencing of the processor
- Provided to improve processor utilization
 - Most I/O devices are slower than the processor
 - Processor must pause to wait for device

Flow of Control w/o Interrupts and I/O Wait

Figure 1

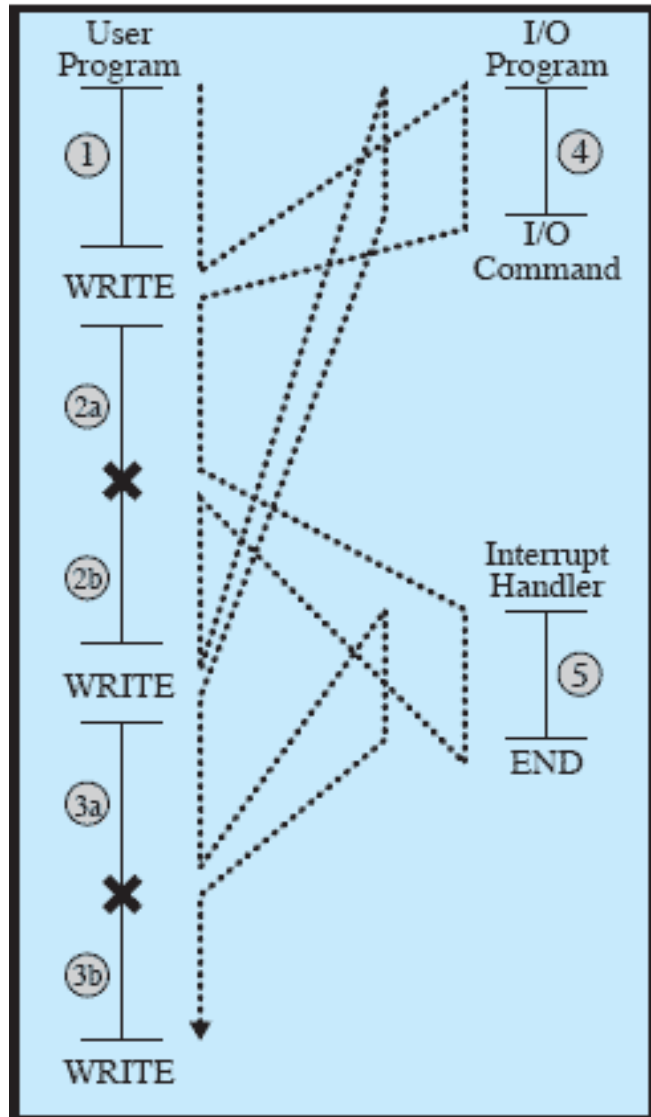


(a) No interrupts




Flow of Control with Interrupts and I/O Wait

Figure 2



(b) Interrupts; short I/O wait



Transfer of Control via Interrupts

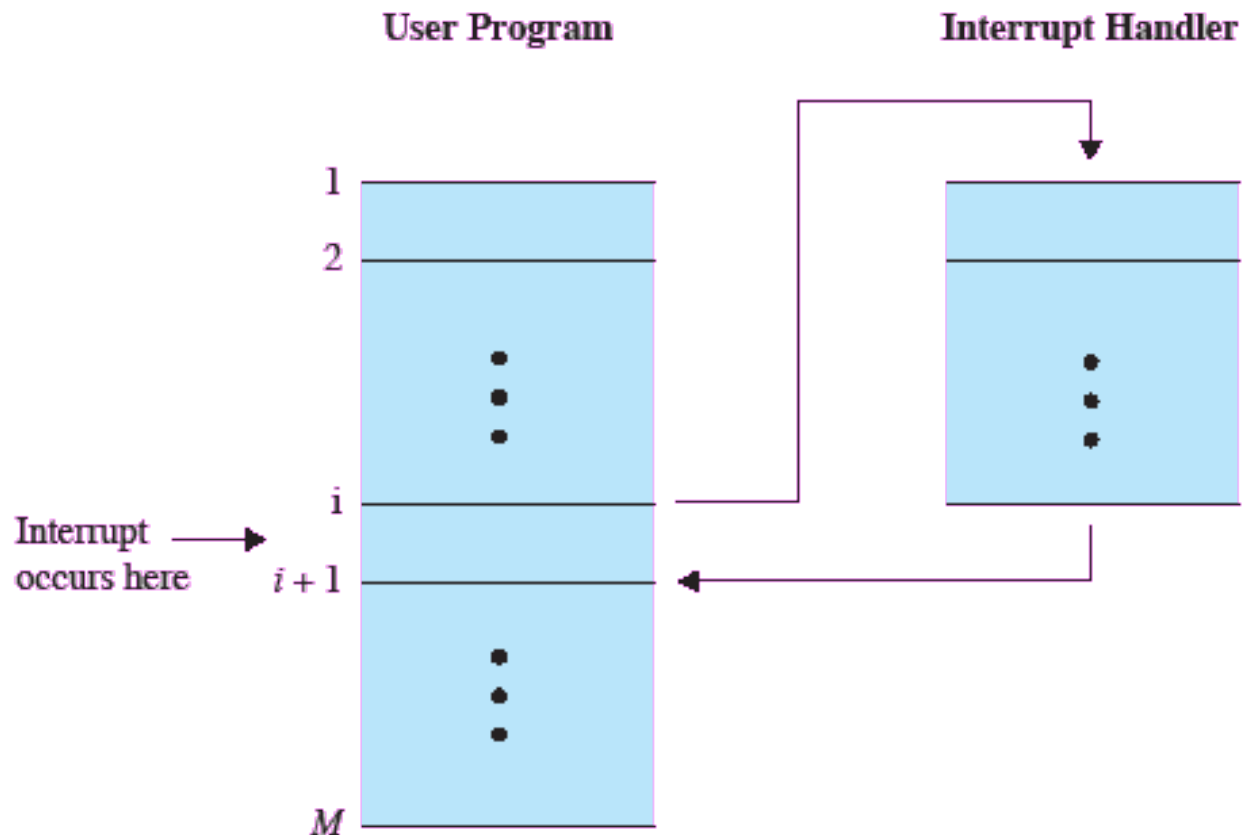


Figure 1.6 Transfer of Control via Interrupts

Instruction Cycle with Interrupts

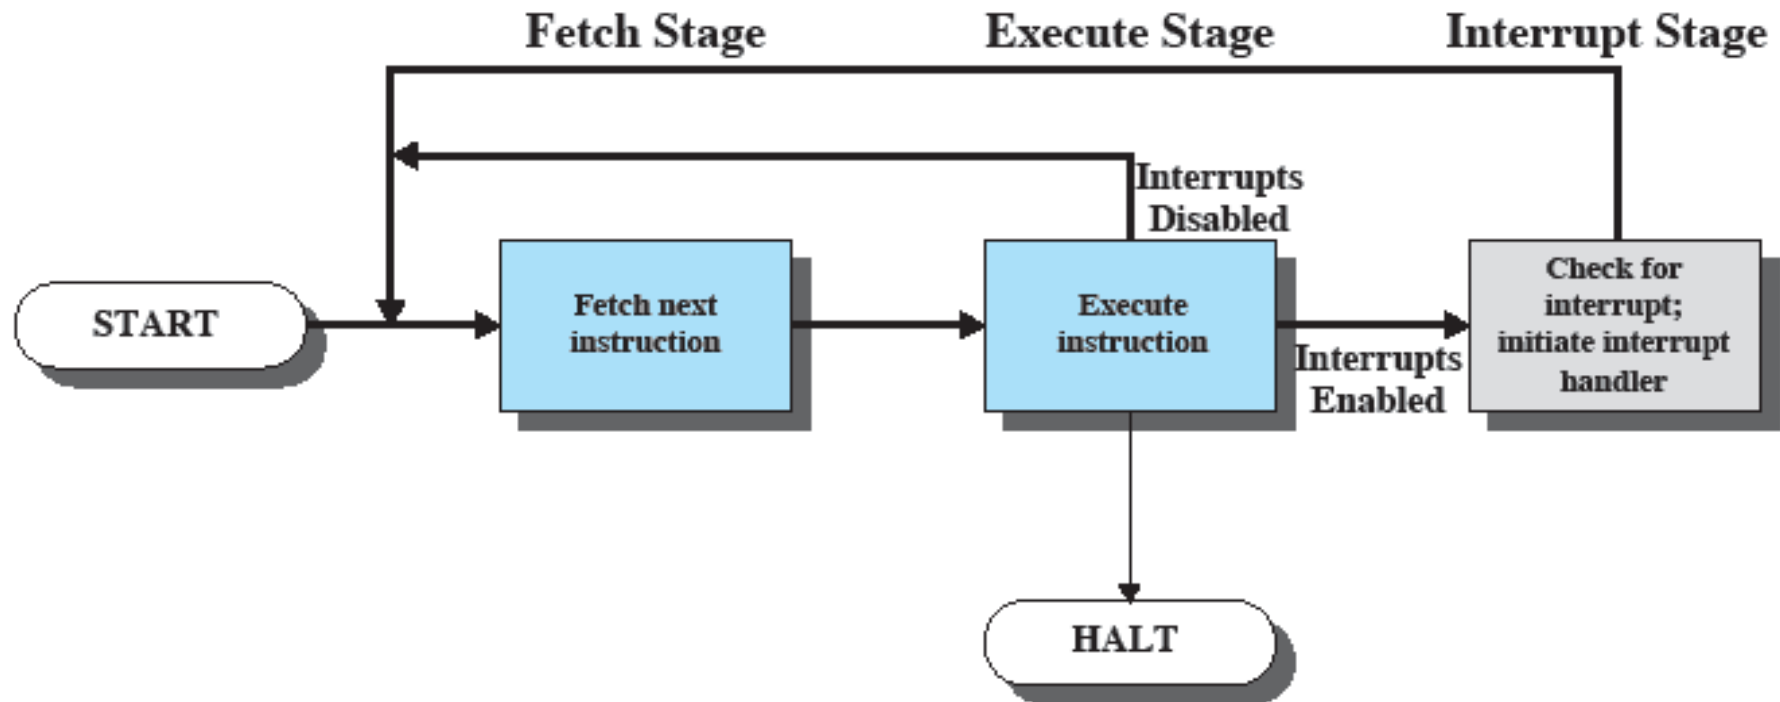


Figure 1.7 Instruction Cycle with Interrupts