


# Inter-Process Communication


# Inter-process Communications

- Passing information between processes
- Used to coordinate process activity
- May result in data inconsistency if mechanisms are not in place to ensure orderly execution of the processes.
- The problem can be identified as a race condition

# Producer / Consumer Problem



# Producer / Consumer Problem



# The Producer Consumer Problem

A **producer** process "produces" information "consumed" by a **consumer** process.

Here are the variables needed to define the problem:

```
#define BUFFER_SIZE 10
typedef struct {
    DATA    data;
} item;
item  buffer[BUFFER_SIZE];
int   in = 0;           // Location of next input to buffer
int   out = 0;          // Location of next removal from buffer
int   counter = 0;      // Number of buffers currently full
```

Consider the code segments on the next slide:


- Does it work?
- Are all buffers utilized?

# PROCESS SYNCHRONIZATION

A **producer** process "produces" information "consumed" by a **consumer** process.

```
item    nextProduced;          PRODUCER

while (TRUE) {
    while (counter == BUFFER_SIZE);
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```



## The Producer Consumer Problem

```
#define BUFFER_SIZE 10
typedef struct {
    DATA    data;
} item;
item    buffer[BUFFER_SIZE];
int     in = 0;
int     out = 0;
int     counter = 0;
```

```
item    nextConsumed;          CONSUMER

while (TRUE) {
    while (counter == 0);
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
}
```

# PROCESS SYNCHRONIZATION

## The Producer Consumer Problem

Note that `counter++;` ← this line is NOT what it seems!!

is really -->

```
register = counter
register = register + 1
counter = register
```

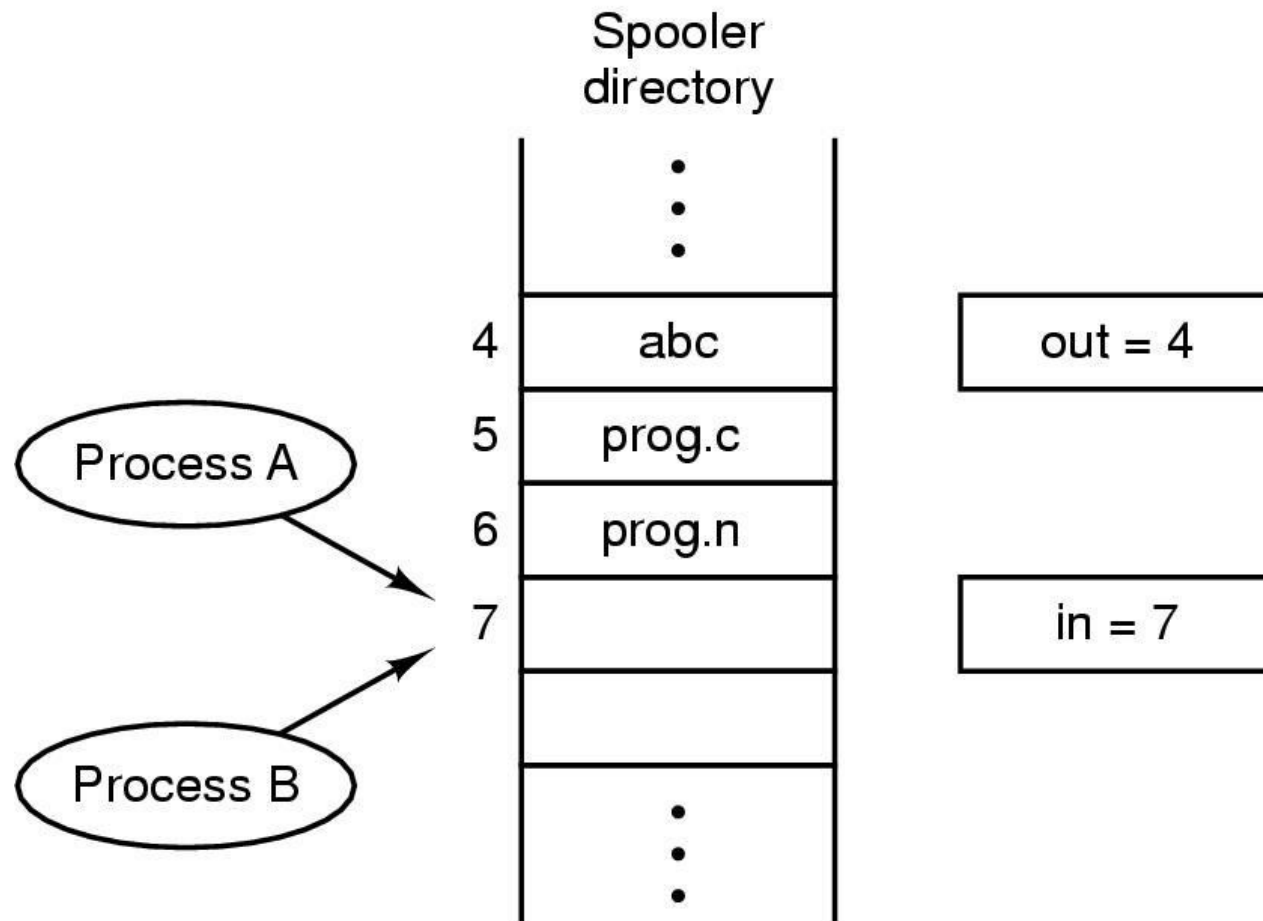
At a micro level, the following scenario could occur using this code:

<b>T0;</b>	<b>Producer</b>	<b>Execute</b>	<b>register1 = counter</b>	<b>register1 = 5</b>
<b>T1;</b>	<b>Producer</b>	<b>Execute</b>	<b>register1 = register1 + 1</b>	<b>register1 = 6</b>
<b>T2;</b>	<b>Consumer</b>	<b>Execute</b>	<b>register2 = counter</b>	<b>register2 = 5</b>
<b>T3;</b>	<b>Consumer</b>	<b>Execute</b>	<b>register2 = register2 - 1</b>	<b>register2 = 4</b>
<b>T4;</b>	<b>Producer</b>	<b>Execute</b>	<b>counter = register1</b>	<b>counter = 6</b>
<b>T5;</b>	<b>Consumer</b>	<b>Execute</b>	<b>counter = register2</b>	<b>counter = 4</b>

**The Problem:**

**atomic execution of counter changes**

# Race Conditions



Two processes want to access shared memory at the same time.



# The Critical Section / Region Problem

- Occurs in systems where multiple processes all compete for the use of shared data.
- Each process includes a section of code (the **critical section**) where it accesses this shared data.
- The problem is to ensure that **only one process at a time is allowed** to be operating in its critical section.

# PROCESS SYNCHRONIZATION

## Critical Section / Region Problem

Critical section must **ENFORCE ALL 3** of the following rules:

**Mutual Exclusion:** No more than one process can execute in its critical section at one time.

**Progress:** If no process is in the critical section and one of them wants to go in, then the interested processes should be able to decide in a finite time who should go in.  
Processes in their remainder section must not block another process interested to enter a critical section.  
Selection of that process cannot be delayed indefinitely.

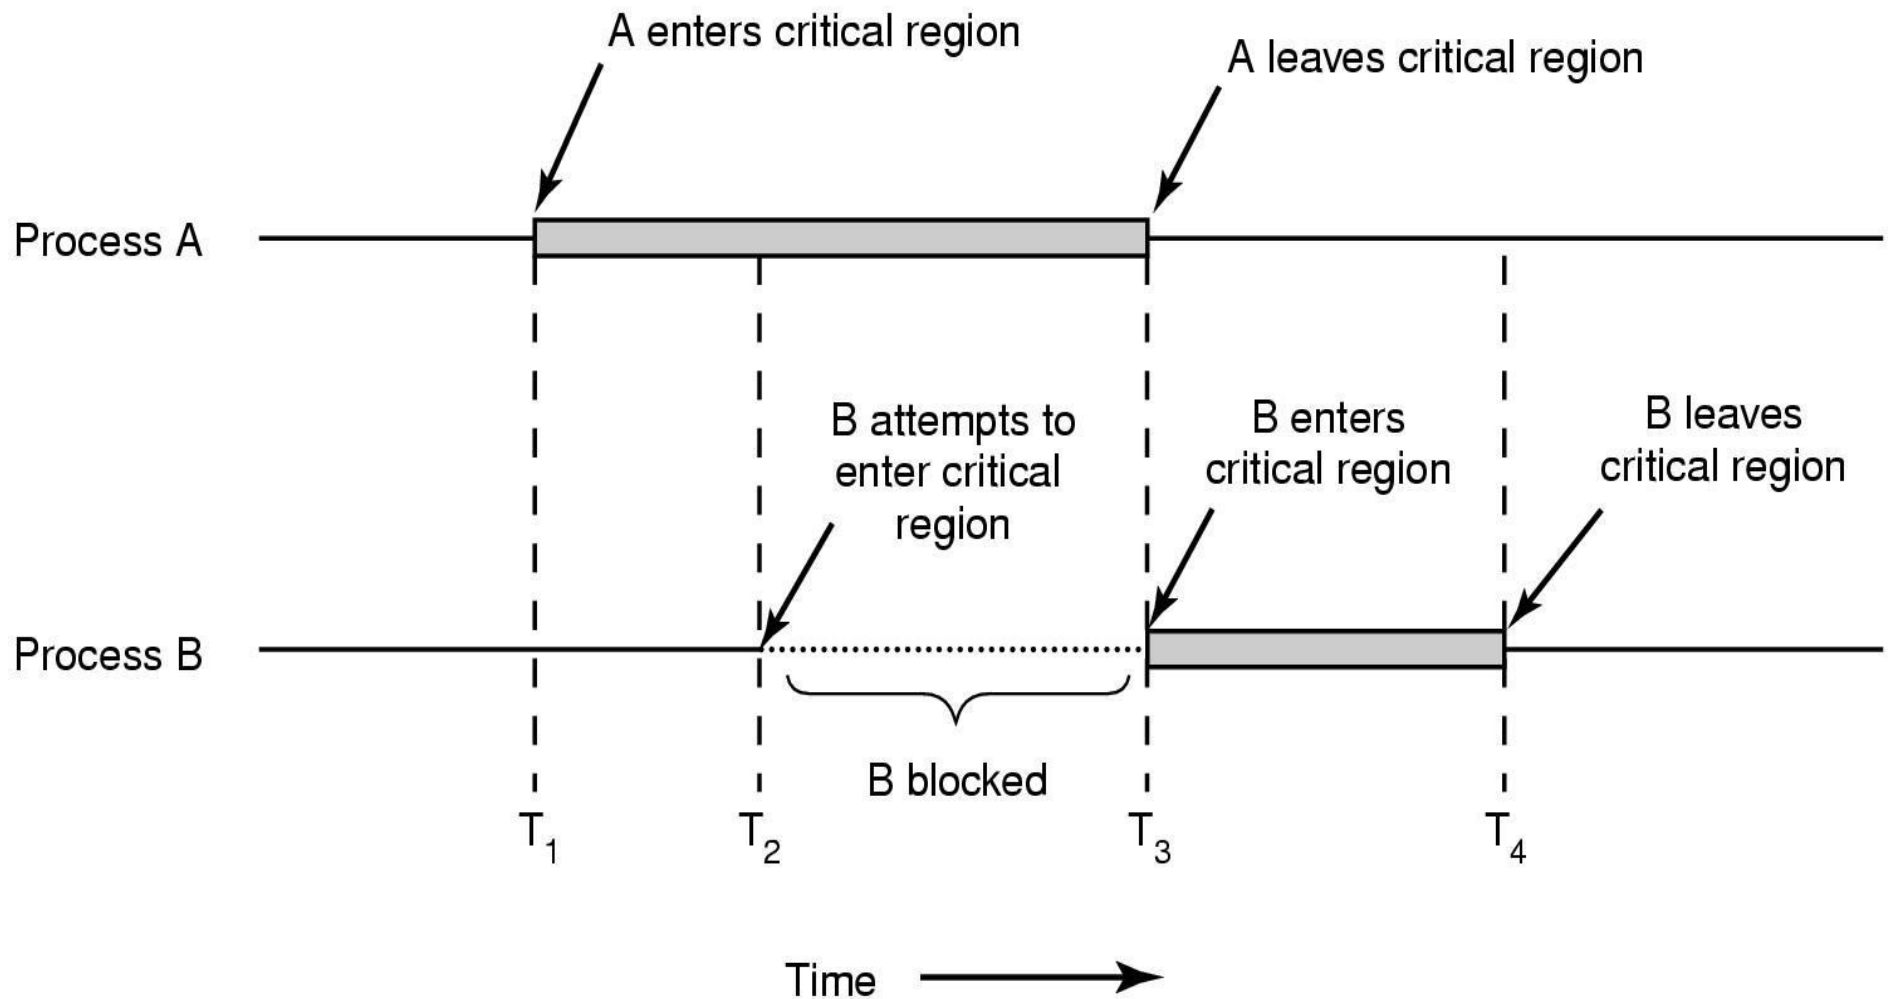
**Bounded Wait:** All requesters (processes) must eventually be let into the critical section.  
There is a bound on the number of times that a waiting process can be superceded.

Can **Race Conditions** be avoided using above conditions?

# Conditions to avoid Race Conditions

- No two processes may be simultaneously inside their critical regions.
- No assumptions may be made about speeds or the number of CPUs.
- No process running outside its critical region may block other processes.
- No process should have to wait forever to enter its critical region.

# How Critical Regions Operate?



**Mutual Exclusion Preserved!**

Are there are other conditions  
we should worry about?

Guess?

# Are there are other conditions we should worry about?

- Starvation
- Deadlock

# Try this Exercise! Lets see your approach

Following solution is alleged to be a solution to the critical section problem.  
Argue for its correctness or show a case in which it fails.

```
shared int turn;  
shared boolean flag[2];
```

```
proc (int i) {  
    while (TRUE) {  
        compute;
```

```
try:  flag[i] = TRUE;
```

```
    while (flag[(i+1) mod 2]) {  
        if (turn == i)  
            continue;  
        flag[i] = FALSE;  
        while (turn != i);  
        goto try;  
    }
```

```
<critical section>;
```

```
    turn = (i+1) mod 2;  
    flag[i] = FALSE;  
    }  
}
```

```
    turn = 0;  
    flag[0] = flag[1] = FALSE;
```

```
    Run proc with arg of 0  
    Run proc with arg of 1
```

Can you extract the structure of the above program?

# Extracted Structure of the Program Code

```
shared int turn;  
shared boolean flag[2];
```

```
proc (int i) {  
  while (TRUE) {  
    compute;
```

```
try:  flag[i] = TRUE;
```

```
  while (flag[(i+1) mod 2]) {  
    if (turn == i)  
      continue;  
    flag[i] = FALSE;  
    while (turn != i);  
    goto try;  
  }
```

```
<critical section>;
```

```
turn = (i+1) mod 2;  
flag[i] = FALSE;  
}
```

```
turn = 0;  
flag[0] = flag[1] = FALSE;
```

```
Run proc with arg of 0  
Run proc with arg of 1
```



# Try the same Exercise (now COMMENTED)!

Following solution is alleged to be a solution to the critical section problem. Argue for its correctness or show a case in which it fails.

shared int turn; // Keeps track of whose turn it is  
shared boolean flag[2]; // If TRUE, indicates that  
// a process would like to enter its c.s.

```
proc (int i) {  
  while (TRUE) {  
    compute;  
    // Attempt to enter the critical section  
    try: flag[i] = TRUE; // An atomic operation  
    // While the other process's flag is TRUE  
    while (flag[(i+1) mod 2]){//atomic operation  
      if (turn == i)  
        continue;  
      flag[i] = FALSE; // Reset to let other  
    }  
    // process go  
    while (turn != i); // Wait till it's my  
    // turn  
    goto try;  
  }  
}
```

// Okay to enter the critical section

<critical section>;

// Leaving the critical section

```
turn = (i+1) mod 2; // Set turn to other process  
flag[i] = FALSE; // Indicate no desire to enter  
// my cs  
}
```

```
turn = 0; // Process 0 wins tie for 1st turn  
flag[0] = flag[1] = FALSE; // Initialize  
// flags before starting
```

Run proc with arg of 0 //process 0

Run proc with arg of 1 //process 1

Visualize how this program will run?

Suppose process 1 is in its "compute;" section for an extremely long period of time, such that process 0 has time to execute its critical section several times. Will process 0 have to wait for process 1? Show why or why not. Can you give a schedule for two processes?

# Solution to Exercise

**The answer is that process 0 will not have to wait for process 1. It can execute its critical section several times while process 1 is in its "compute;" section.**

# Solution to Exercise

**The answer is that process 0 will not have to wait for process 1. It can execute its critical section several times while process 1 is in its "compute;" section.**

Assume turn is 0 to begin with. If process 1 is in its compute section, we know that flag[1] is FALSE. Process 0 attempts to enter its critical section, so it sets flag[0] to TRUE. Since flag[1] is FALSE, process 0 doesn't enter the while loop. It executes its critical section, and then sets turn to 1 and flag[0] to FALSE.

Process 0 then returns to the top of the while(TRUE) loop and computes. Then it attempts to enter its critical section (Process 1 is still computing). Process 0 sets flag[0] to TRUE, and again it does not enter the while loop, but moves directly to its critical section code.

Although the turn variable was set to 1 (which indicates that it is the other process's turn, it did not enter the while loop, so it didn't have to test turn==i (which would have failed).

# PROCESS SYNCHRONIZATION

## Critical Regions

**A section of code, common to  $N$  cooperating processes, in which the processes may be accessing common variables.**

A Critical Section Environment contains:

<b>Entry Section</b>	Code requesting entry into the critical section.
<b>Critical Section</b>	Code in which only one process can execute at any one time.
<b>Exit Section</b>	The end of the critical section, releasing or allowing others in.
<b>Remainder Section</b>	Rest of the code AFTER the critical section.

# Few Approaches to Achieve Mutual Exclusion

Two categories of solutions –

- ***Hardware solution*** - Disabling (DI / EI) interrupts, special instructions like TSL / CMP / SWAP
- ***Software solution*** - Strict alternation, Lock variables, Peterson's solution and few more.

# PROCESS SYNCHRONIZATION

## Two Process Software

Here's an example of a simple piece of code containing the components required in a critical section.

```
do {  
  /* check conditions for CS */  
  /* critical section */  
  /* set conditions for CS */  
  /* remainder section */  
} while(TRUE);
```

**Entry Section**

**Critical Section (CS)**

**Exit Section**

**Remainder Section**

We will now try a succession of increasingly complicated solutions to the problem of creating valid entry / exit sections.

# Approach 1:

Using one variable for process turns

# Strict Alternation

**NOTE:** In all examples,  $i$  is the current process,  $j$  the "other" process. In these examples, envision the same code running for two processes at the same time.

## TOGGLED ACCESS:

```
do {  
  while ( turn  $\neq$  i );  
  /* critical section */  
  turn = j;  
  /* remainder section */  
} while(TRUE);
```

### Algorithm 1

Are the three Critical Section Requirements Met?

Also called Strict Alternation

A proposed solution to the critical region problem. (a) Process 0. (b) Process 1. In both cases, be sure to note the semicolons terminating the while statements.

```
while (TRUE) {  
  while (turn  $\neq$  0)      /* loop */;  
  critical_region();  
  turn = 1;  
  noncritical_region();  
}
```

(a)

```
while (TRUE) {  
  while (turn  $\neq$  1)      /* loop */;  
  critical_region();  
  turn = 0;  
  noncritical_region();  
}
```

(b)



Approach 2:  
Using one variable indicating  
interest to enter CR

# Lock Variables

## FLAG FOR EACH PROCESS GIVES STATE:

Each **process** maintains a **flag** indicating that it **wants to get into the critical section**.

It **checks the flag of other process** and doesn't enter the critical section if that other process wants to get in.

### Shared variables

☞ **boolean** flag[2];

initially **flag [0] = flag [1] = false**.

☞ **flag [i] = true**  $\Rightarrow P_i$  ready to enter its critical section

```
do {  
  flag[i] := true;  
  while (flag[j]) do no-op;          /*  
  critical section */  
  flag [i] = false;  
  /* remainder section */  
} while (1);
```

**Algorithm 2**

**Are the three Critical  
Section Requirements Met?**

**Principle:** **Set** the flag (LOCK) to enter CR, and **Reset** the flag after exit

# Approach 3:

## Using turns and locks together

# PROCESS SYNCHRONIZATION

## Two Processes Software

### FLAG TO REQUEST ENTRY:

- Each processes sets a flag to request entry. Then each process toggles a bit to allow the other in first.
- This code is executed for each process  $i$ .

### Shared variables

☞ **boolean** `flag[2]`;

initially `flag [0] = flag [1] = false`.

☞ `flag [i] = true`  $\Rightarrow P_i$  ready to enter its critical section

```
do {  
  flag [i] := true;  
  turn = j;  
  while (flag [j] and turn == j);  
  /* critical section */  
  flag [i] = false;  
  /* remainder section */  
} while (1);
```

Algorithm 3

Are the three Critical Section  
Requirements Met?

This is Peterson's  
Solution

How this is different then  
previous solutions?

# Peterson's Solution to achieve Mutual Exclusion

```
#define FALSE 0
#define TRUE 1
#define N      2                /* number of processes */

int turn;                       /* whose turn is it? */
int interested[N];              /* all values initially 0 (FALSE) */

void enter_region(int process);  /* process is 0 or 1 */
{
    int other;                  /* number of the other process */

    other = 1 - process;        /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;             /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process)   /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

# Producer Consumer Problem

- Using strict alternation
- Using lock variables
- Using Peterson's solution

Write pseudocode and verify for its correctness.

# PROCESS SYNCHRONIZATION

## Critical Sections

The hardware required to support critical sections must have (minimally):

- Indivisible instructions (what are they?)
- Atomic load, store, test instruction.
  - For instance, if a store and test occur simultaneously, the test gets EITHER the old or the new, but not some combination.
    - **Atomic = non-interruptable**
- Two atomic instructions, if executed simultaneously, behave as if executed sequentially.

# PROCESS SYNCHRONIZATION

## Hardware Solutions

### Disabling Interrupts:

Works for the Uniprocessor case only.

Needs a modified approach for multiprocessor / multi-core processors.

**Disable interrupts** (for e.g., DI)

**/\* critical region \*/**

**Enable interrupts** (for e.g., EI)

### Atomic test and set (Use of TSL instruction)

Returns parameter & sets parameter to true atomically.

**while ( test\_and\_set ( lock ));**

**/\* critical section \*/**

**lock = false;**



# The TSL Instruction ...(1)

enter\_region:

TSL REGISTER,LOCK

CMP REGISTER,#0

JNE enter\_region

RET

| copy lock to register and set lock to 1

| was lock zero?

| if it was nonzero, lock was set, so loop

| return to caller; critical region entered

leave\_region:

MOVE LOCK,#0

RET

| store a 0 in lock

| return to caller

**Entering and leaving a critical region using the TSL instruction.**

# The XCHG Instruction ...(2)

enter\_region:

MOVE REGISTER,#1

| put a 1 in the register

XCHG REGISTER,LOCK

| swap the contents of the register and lock variable

CMP REGISTER,#0

| was lock zero?

JNE enter\_region

| if it was non zero, lock was set, so loop

RET

| return to caller; critical region entered

leave\_region:

MOVE LOCK,#0

| store a 0 in lock

RET

| return to caller

**Entering and leaving a critical region using the XCHG instruction.**

# Can we find a solution to busy waiting?

Can we have a mechanism where a process is not constantly checking for the availability of CR, rather is being informed about the availability of CR as and when that scenario arises?

# Can we find a solution to busy waiting?

Can we have a mechanism where a process is not constantly checking for the availability of CR, rather is being informed about the availability of CR as and when that scenario arises?

## SLEEP and WAKEUP operations

A call to **SLEEP** blocks the calling process.

A call to **WAKEUP** unblocks a process that is passed as an argument in the call.

# Producer-Consumer Problem

Assume there are two special operations – sleep and wakeup.

**Write a pseudocode for the producer-consumer problem using these above operations.**

**Analyze your pseudocode.**

# Pseudocode for Producer-Consumer Problem

```
#define N 100
int count = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item( );
        if (count == N) sleep( );
        insert_item(item);
        count = count + 1;
        if (count == 1) wakeup(consumer);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep( );
        item = remove_item( );
        count = count - 1;
        if (count == N - 1) wakeup(producer);
        consume_item(item);
    }
}
```

/\* number of slots in the buffer \*/  
/\* number of items in the buffer \*/

/\* repeat forever \*/  
/\* generate next item \*/  
/\* if buffer is full, go to sleep \*/  
/\* put item in buffer \*/  
/\* increment count of items in buffer \*/  
/\* was buffer empty? \*/

/\* repeat forever \*/  
/\* if buffer is empty, got to sleep \*/  
/\* take item out of buffer \*/  
/\* decrement count of items in buffer \*/  
/\* was buffer full? \*/  
/\* print item \*/

Are there any issues with the above code?